Persistent Memory Hackathon and Workshop Flash Memory World 2019

August 23, 2019

https://github.com/pmemhackathon/2019-08-23

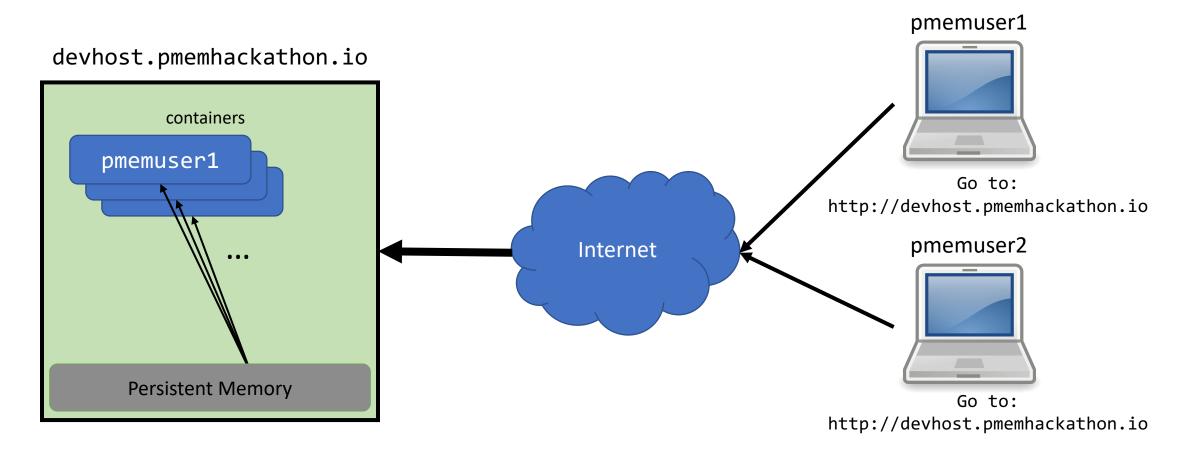


Agenda

- Essential Background Slides, covering:
 - Logistics: how you access persistent memory from your laptop
 - The minimum you need to know about persistent memory
 - Walk through the first example or two together
- Less talk, more hack...
 - Work your way through the examples, in any order after the first three
 - Helpers will be available in the room
 - If FAQs come up, we'll present answers to the entire room



Logistics: The webhackathon Tool



• • •





Webhackathon Basics

- List of examples presented on main page
 - First three recommended to provide essential background
 - We will walk through some of these together
 - Pick examples that are interesting to you (task, language, etc)
 - Use them as a starting point for your own code
- Menu provides:
 - Access to these background slides
 - Browse your copy of the repo (to download something you want to keep)
 - Browser-based shell window for your container (for users who need it)
- Everything you do runs in your own container on the server
 - With your own copy of the hackathon repo
 - The path to the persistent memory is /pmem
- We're all friends here: please no denial-of-service attacks on server!

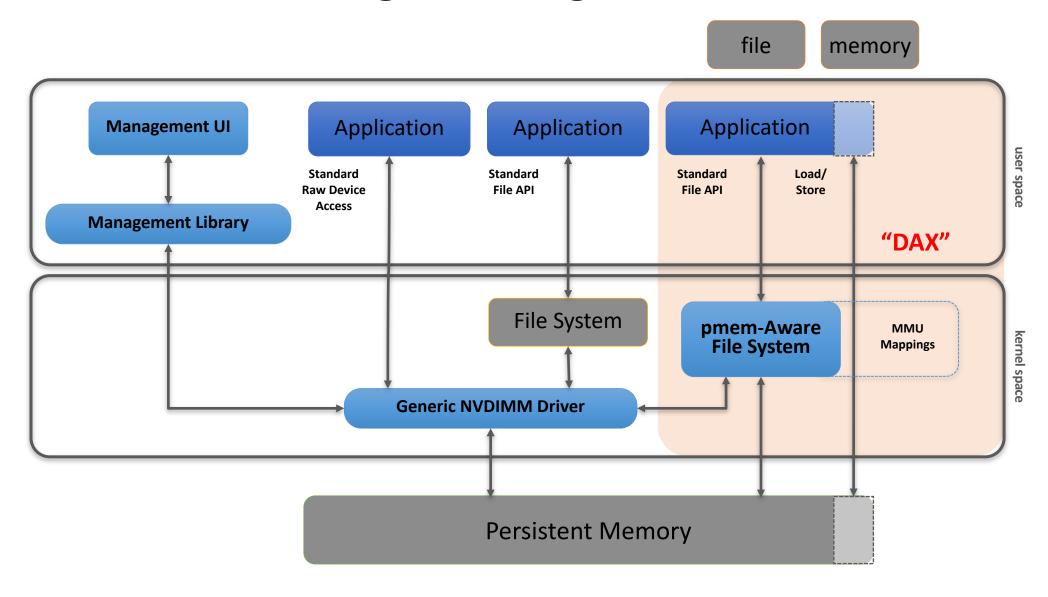


Essential pmem Programming Background

- Lots of ways to use pmem with existing programs
 - Storage APIs
 - Libraries or kernels using pmem transparently
 - Memory Mode
- This hackathon doesn't cover the above (too easy!)
 - We assume you want direct access to pmem
 - We show code, but also concepts
 - There are lots of paths you can take, these are just examples

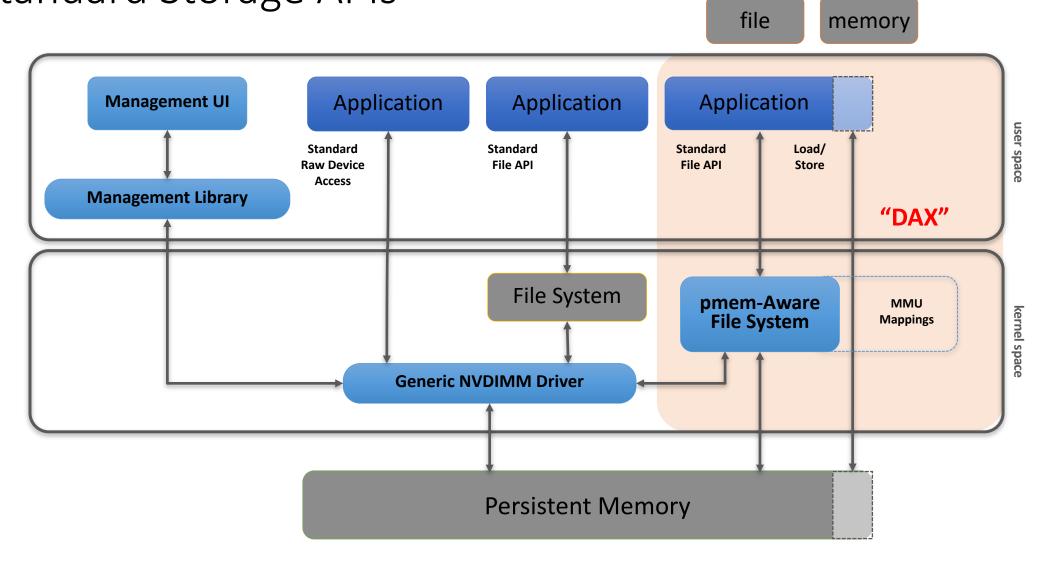


The SNIA NVM Programming Model

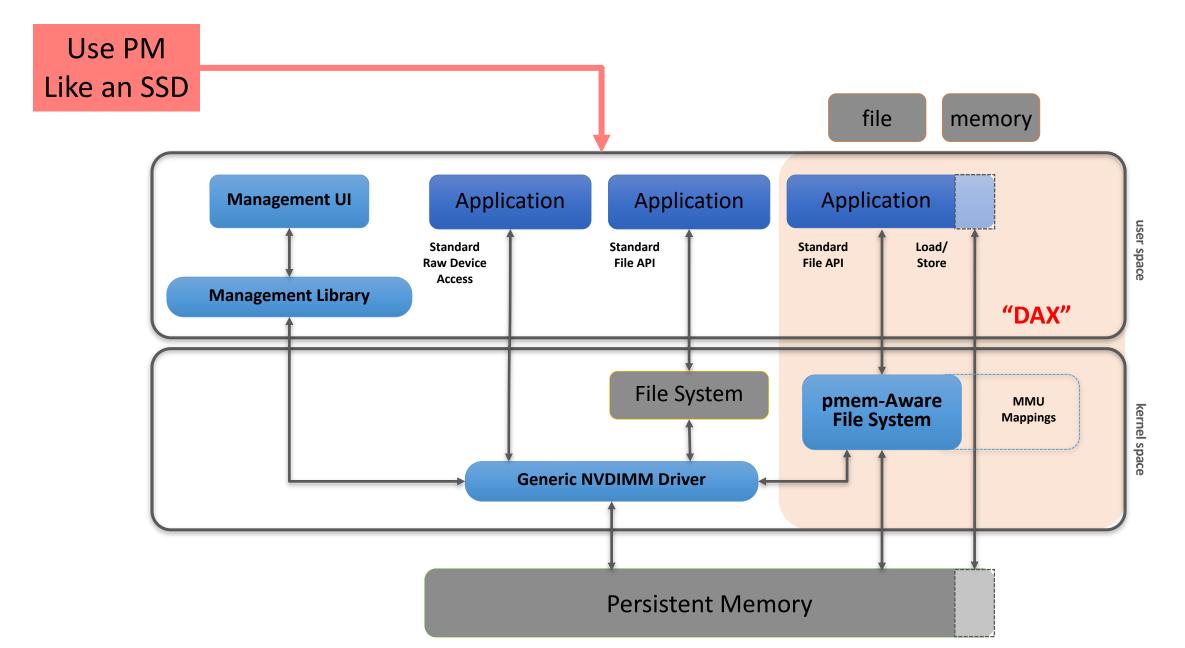




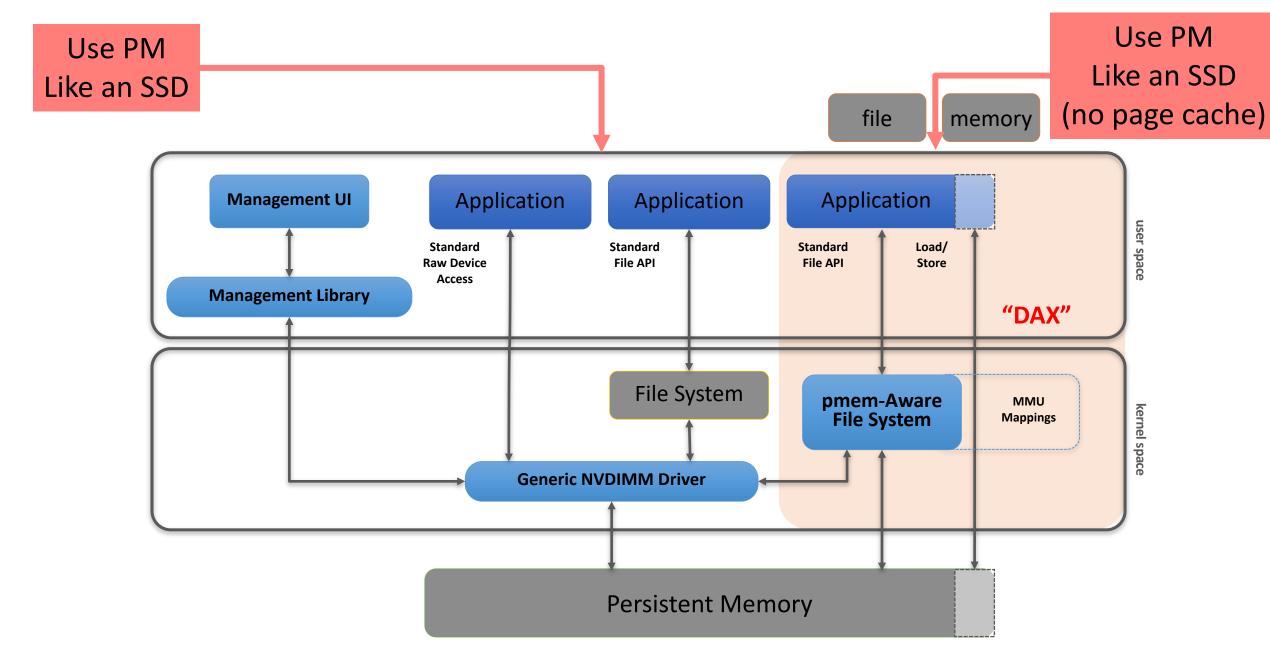
Don't Forget: The NVM Programming Model Starts With Standard Storage APIs



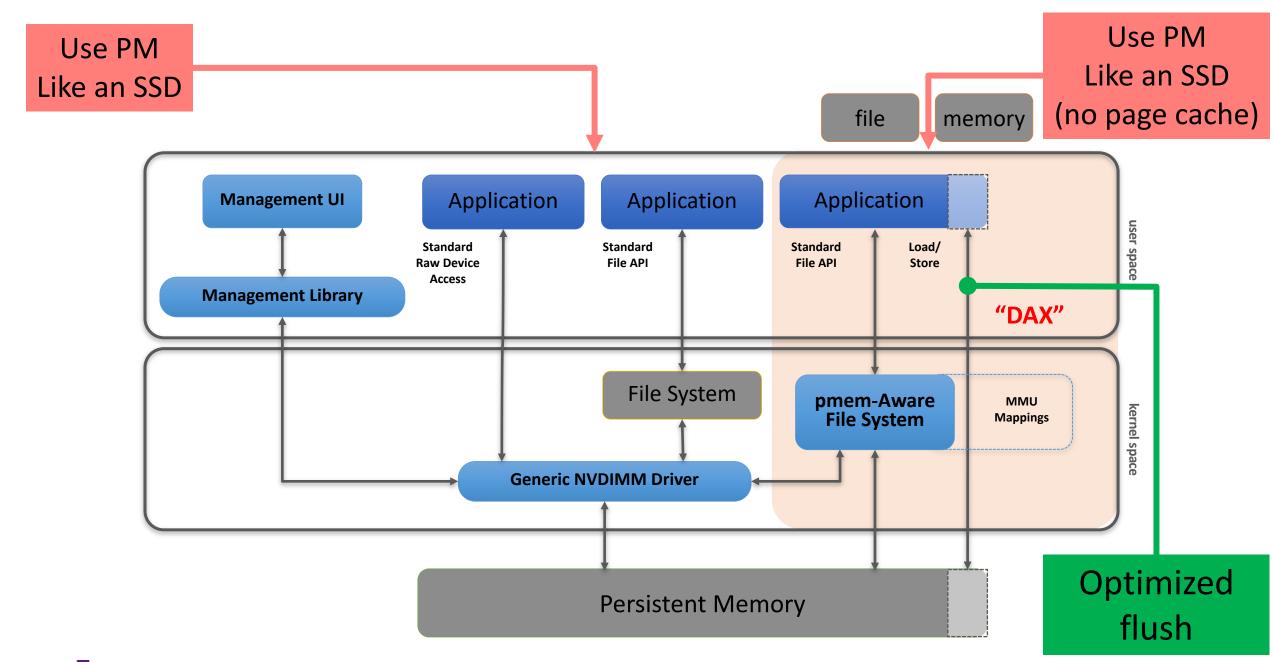






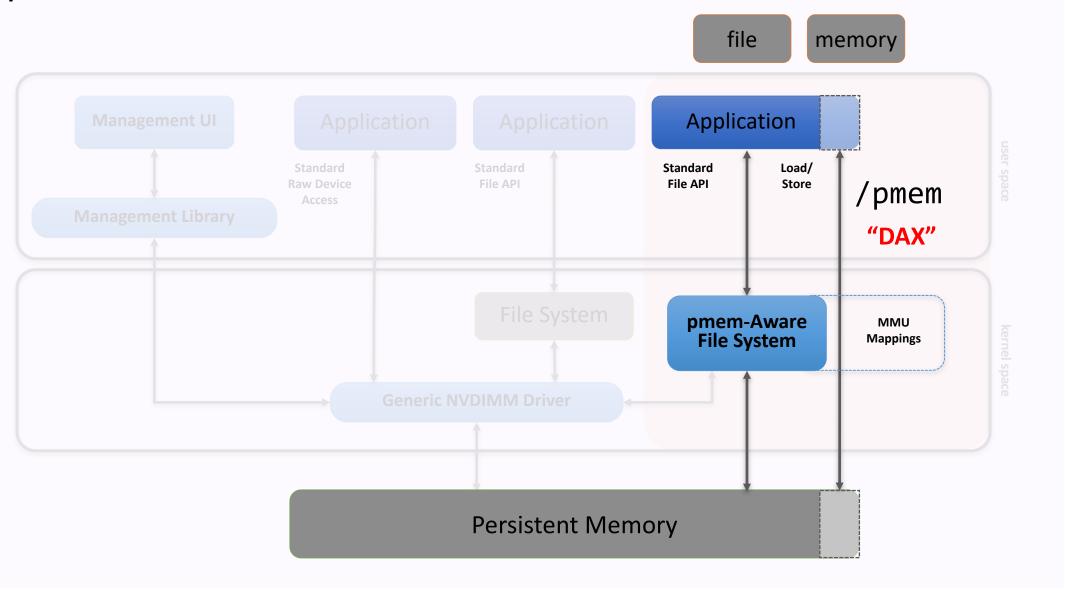








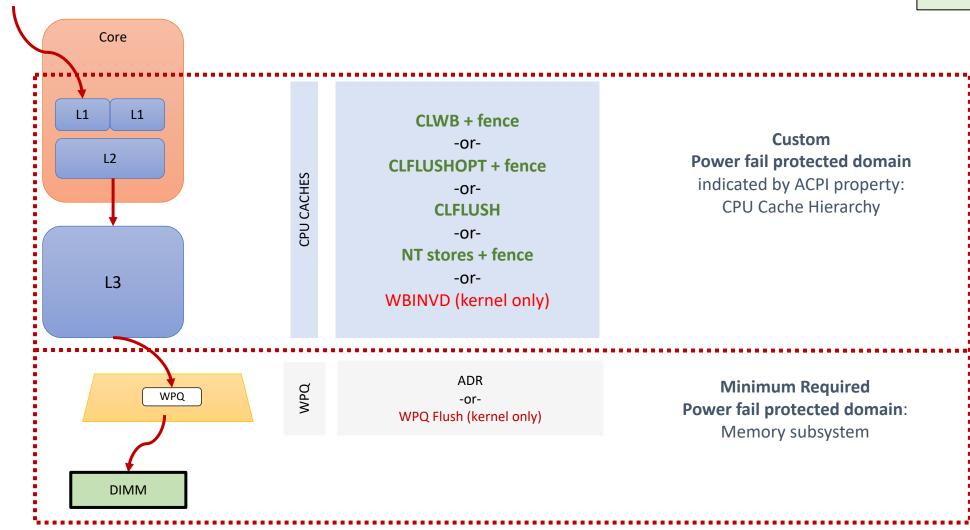
Today's focus





How the Hardware Works

Not shown: MCA ADR Failure Detection

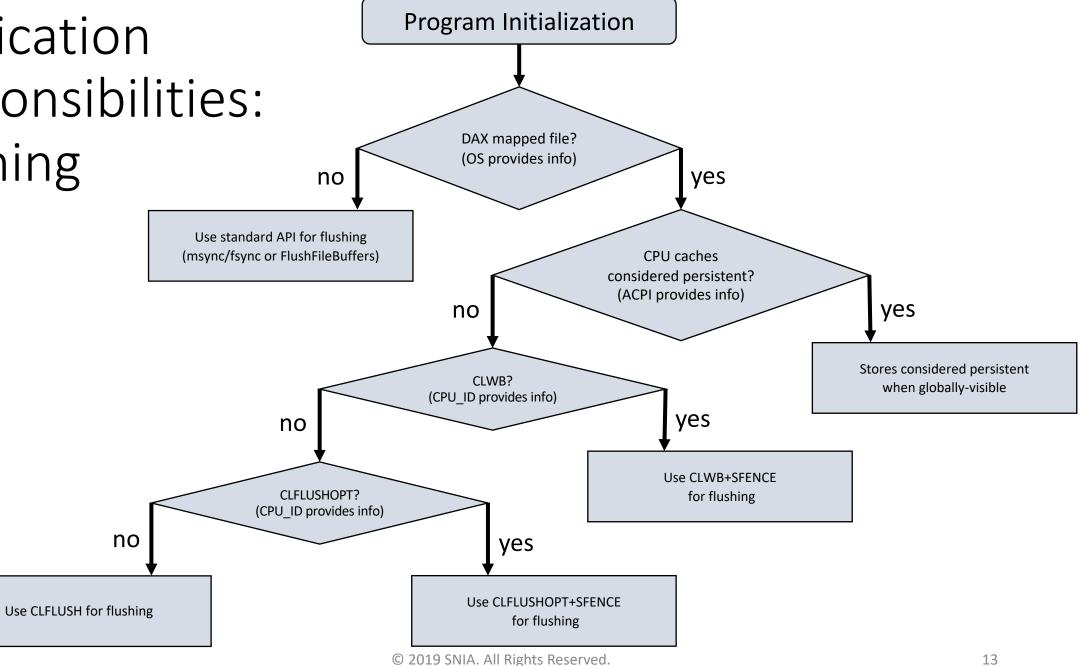




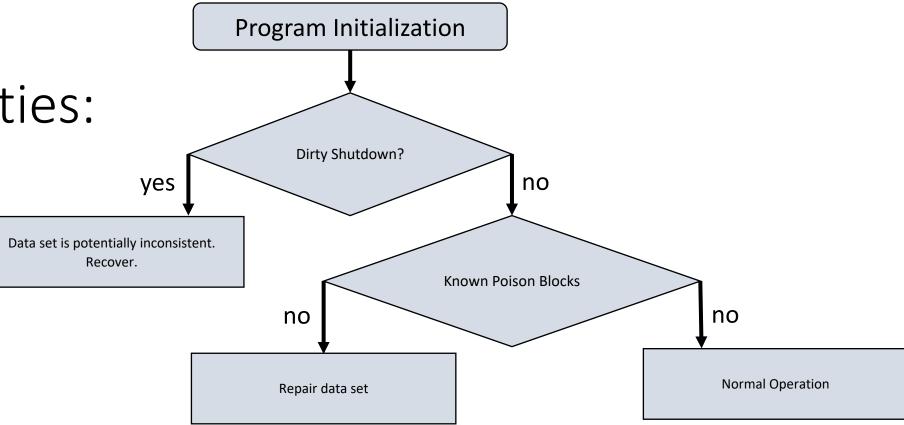
MOV

Application Responsibilities: Flushing

no



Application Responsibilities: Recovery





Application Responsibilities: Consistency

```
open(...);

mmap(...);

strcpy(pmem, "Hello, World!");

msync(...);
Crash
```

Result

```
    "\0\0\0\0\0\0\0\0\0\0\0..."
    "Hello, W\0\0\0\0\0\0..."
    "\0\0\0\0\0\0\0\0\0\0\0\0"
    "Hello, \0\0\0\0\0\0\0\0\0"
    "Hello, World!\0"
```



Application Responsibilities: Consistency

```
open(...);
mmap(...);
strcpy(pmem, "Hello, World!");
pmem_persist(pmem, 14);
Crash
```

pmem_persist() may be faster,
but is still not transactional

Result

"\0\0\0\0\0\0\0\0\0\0\0..."
 "Hello, W\0\0\0\0\0\0..."
 "\0\0\0\0\0\0\0\0\0\0\0\0"
 "Hello, \0\0\0\0\0\0\0\0\0"
 "Hello, World!\0"



The Persistent Memory Development Kit PMDK http://pmem.io

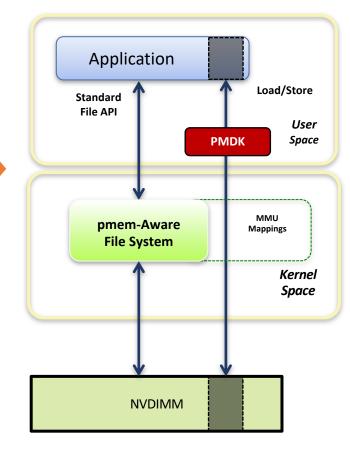
- PMDK is a collection of libraries
 - Developers pull only what they need
 - Low level programming support
 - Transaction APIs
 - Fully validated
 - Performance tuned.
- Open Source & Product neutral





PMDK Libraries

PCJ – Persistent C++ **Pvthon** pmemkv Collection for Java Interface for persistent memory Interface to create a Interface to create arrays of allocation, transactions and persistent memory Transaction pmem-resident blocks, of resident log file general facilities Support same size, atomically updated libpmemlog libpmemobj libpmemblk



Support for volatile memory usage

Low level support for local persistent memory

Low level support for remote access to persistent memory

librpmem

Low-level support

In Development



Hack, hack, hack...

http://devhost.pmemhackathon.io

Username: pmemuserX (handed out)

Password: (handed out)



More Background Information

Read as necessary, or just keep working through the examples – whatever works best for you



Resources

- PMDK Resources:
 - Home: https://pmem.io
 - PMDK: https://pmem.io/pmdk
 - PMDK Source Code : https://github.com/pmem/PMDK
 - Google Group: https://groups.google.com/forum/#!forum/pmem
 - Intel Developer Zone: https://software.intel.com/persistent-memory
 - Memkind: https://github.com/memkind/memkind (see memkind_pmem(3))
 - libpmemkv: https://github.com/pmem/pmemkv
- NDCTL: https://pmem.io/ndctl
- SNIA NVM Programming Model: https://www.snia.org/tech_activities/standards/curr_standards/npm
- Getting Started Guides: https://docs.pmem.io



More Developer Resources

- Find the PMDK (Persistent Memory Development Kit) at http://pmem.io/pmdk/
- Getting Started
 - Intel IDZ persistent memory- https://software.intel.com/en-us/persistent-memory
 - Entry into overall architecture http://pmem.io/2014/08/27/crawl-walk-run.html
 - Emulate persistent memory http://pmem.io/2016/02/22/pm-emulation.html
- Linux Resources
 - Linux Community Pmem Wiki https://nvdimm.wiki.kernel.org/
 - Pmem enabling in SUSE Linux Enterprise 12 SP2 https://www.suse.com/communities/blog/nvdimm-enabling-suse-linux-enterprise-12-service-pack-2/
- Windows Resources
 - Using Byte-Addressable Storage in Windows Server 2016 https://channel9.msdn.com/Events/Build/2016/P470
 - Accelerating SQL Server 2016 using Pmem https://channel9.msdn.com/Shows/Data-Exposed/SQL-Server-2016-SCM--FAST
- Other Resources
 - SNIA Persistent Memory Summit 2018 https://www.snia.org/pm-summit
 - Intel manageability tools for Pmem https://01.org/ixpdimm-sw/

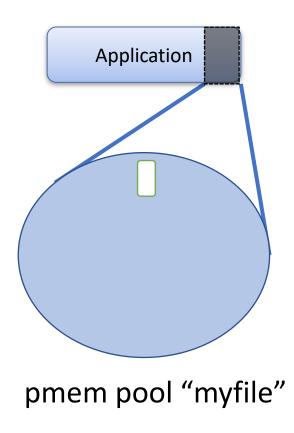


Basic libpmemobj Information

This is the most flexible of the PMDK libraries, supporting general-purpose allocation & transactions



The Root Object

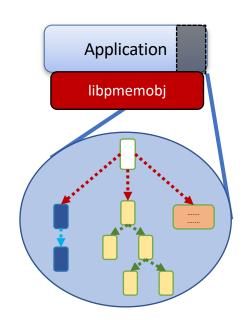


root object:

- assume it is always there
- created first time accessed
- initially zeroed

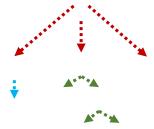


Using the Root Object



Link pmem data structures in pool off the root object to find them on each program run

"pointers" are really Object IDs





C Programming with libpmemobj



Transaction Syntax

```
TX_BEGIN(Pop) {
                  /* the actual transaction code goes here... */
} TX_ONCOMMIT {
                   * optional - executed only if the above block
                   * successfully completes
                   */
} TX_ONABORT {
                  /*
                   * optional - executed if starting the transaction fails
                   * or if transaction is aborted by an error or a call to
                   * pmemobj tx abort()
                   */
} TX_FINALLY {
                   * optional - if exists, it is executed after
                   * TX_ONCOMMIT or TX_ONABORT block
                   */
} TX_END /* mandatory */
```



Properties of Transactions

```
Powerfail
                                               Multi-Thread
           Atomicity
                                                 Atomicity
TX_BEGIN_PARAM(Pop, TX_PARAM_MUTEX, &D_RW(ep)->mtx, TX_PARAM_NONE) {
       TX_ADD(ep);
       D_RW(ep)->count++;
} TX_END
```

Caller must instrument code for undo logging



C++ Programming with libpmemobj



C++ Queue Example: Declarations

```
/* entry in the queue */
struct pmem_entry {
    persistent_ptr<pmem_entry> next;
    p<uint64_t> value;
};
```

persistent_ptr <t></t>	Pointer is really a position-independent Object ID in pmem. Gets rid of need to use C macros like D_RW()
p< <i>T</i> >	Field is pmem-resident and needs to be maintained persistently. Gets rid of need to use C macros like TX_ADD()



C++ Queue Example: Transaction

```
void push(pool base &pop, uint64 t value) {
     transaction::run(pop, [&] {
          auto n = make persistent<pmem entry>();
          n->value = value;
          n->next = nullptr;
          if (head == nullptr) {
               head = tail = n;
          } else {
                                           Transactional
               tail->next = n;
                                      (including allocations & frees)
               tail = n;
```



Intel Developer Support & Tools

PMDK Tools

- Valgrind plugin: pmemcheck
- Debug mode, tracing, pmembench, pmreorder

pmem.io

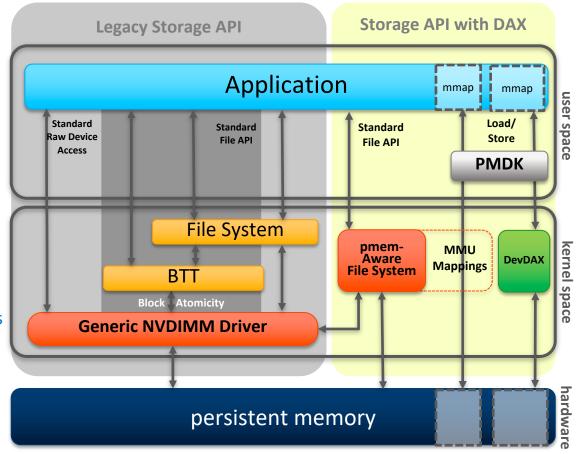
- New features to support Intel[®] Optane[™] DC persistent memory
 - Intel® VTune™ Amplifier Performance Analysis
 - Intel® Inspector Persistence Inspector finds missing cache flushes & more
 - Free downloads available

software.intel.com/pmem



Possible ways to access persistent memory

- No Code Changes Required
- Operates in Blocks like SSD/HDD
 - Traditional read/write
 - Works with Existing File Systems
 - Atomicity at block level
 - Block size configurable
 - 4K, 512B*
- NVDIMM Driver required
 - Support starting Kernel 4.2
- Configured as Boot Device
- Higher Endurance than Enterprise SSDs
- High Performance Block Storage
 - Low Latency, higher BW, High IOPs



- Code changes may be required*
- Bypasses file system page cache
- Requires DAX enabled file system
 - XFS, EXT4, NTFS
- No Kernel Code or interrupts
- No interrupts
- Fastest IO path possible

^{*}Requires Linux



^{*} Code changes required for load/store direct access if the application does not already support this.

Hackathon Contributors...

- Piotr Balcer
- Eduardo Berrocal
- Jim Fister
- Stephen Bates
- Zhiming Li
- Lukasz Plewa

- Szymon Romik
- Andy Rudoff
- Steve Scargall
- Peifeng Si
- Pawel Skowron
- Usha Upadhyayula

With lots of input & feedback from others along the way...

